

## Work Distribution for baseline RAG flow

Cyril - Chunking + Reciprocal Rank Fusion

Use the Unstructured library to create chunks and tune for best chunk size. Implement Reciprocal Rank Fusion to fuse the keyword based and semantic based results.

Jeslin - Embeddings + Vector DB

Use Sentence Transformers to implement an embedding function. Experiment with the models available in the library. Implement search using FAISS as the Vector DB

Nazil - Table handling + Prompt Engineering

Use Camelot to implement table handling. Craft the LLM prompt for answer generation; with a balance between answer quality and hallucination.

Abin - BM25 Keyword Matching + Function creation

Implement BM25 Matching using the rank\_bm25 library and tune the hyperparameters. Turn the chunking, embedding, search, answer- generation components into standalone functions.